


Java Threads

Multitasking and Multithreading

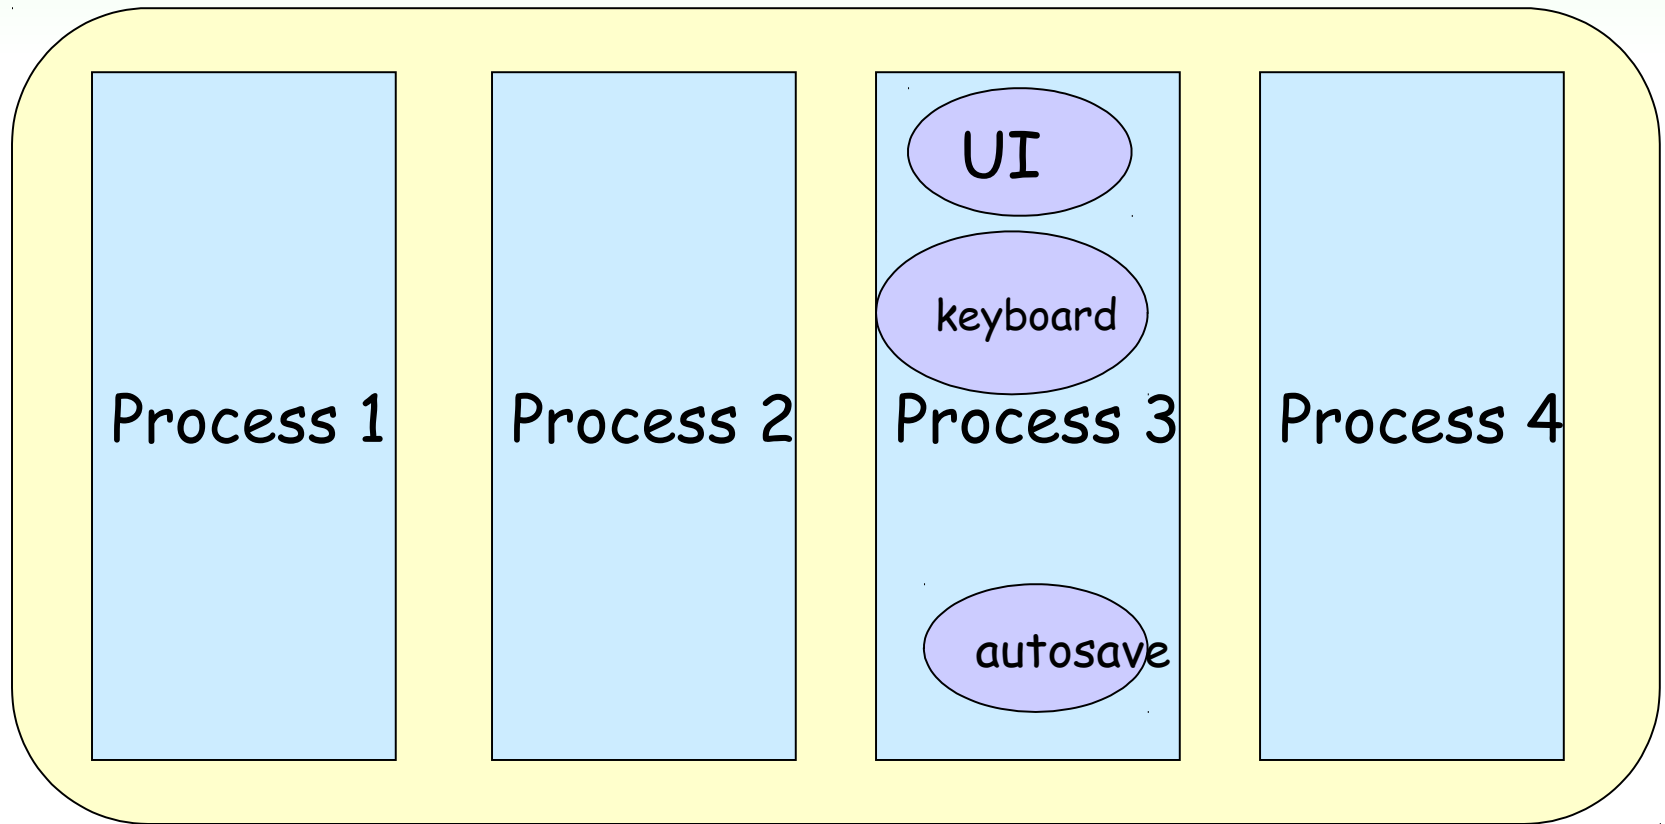
- Multitasking refers to a computer's ability to perform multiple jobs concurrently
 - more than one program are running concurrently, e.g., UNIX
- A thread is a single sequence of execution within a program
- Multithreading refers to multiple threads of control within a single program
 - each program can run multiple threads of control within it, e.g., Web Browser

Concurrency vs. Parallelism



Threads and Processes

CPU



What are Threads Good For?

- To maintain responsiveness of an application during a long running task.
- To enable cancellation of separable tasks.
- Some problems are intrinsically parallel.
- To monitor status of some resource (DB).
- Some APIs and systems demand it: Swing.

Application Thread

- When we execute an application:
 - The JVM creates a Thread object whose task is defined by the **main()** method
 - It starts the thread
 - The thread executes the statements of the program one by one until the method returns and the thread dies

Multiple Threads in an Application

- Each thread has its private run-time stack
- If two threads execute the same method, each will have its own copy of the local variables the methods uses
- However, all threads see the same dynamic memory (heap)
- Two different threads can act on the same object and same static fields concurrently

Thread Methods

void start()

- Creates a new thread and makes it runnable
- This method can be called only once

void run()

- The new thread begins its life inside this method

void stop() (deprecated)

- The thread is being terminated

Thread Methods

- **yield()**
 - Causes the currently executing thread object to temporarily pause and allow other threads to execute
 - Allow only threads of the same priority to run
- **sleep(int *m*)/sleep(int *m*,int *n*)**
 - The thread sleeps for *m* milliseconds, plus *n* nanoseconds

Creating Threads

- There are two ways to create our own **Thread** object
 1. Subclassing the **Thread** class and instantiating a new object of that class
 2. Implementing the **Runnable** interface
- In both cases the **run()** method should be implemented

Extending Thread

```
public class ThreadExample extends Thread {  
    public void run () {  
        for (int i = 1; i <= 100; i++) {  
            System.out.println("Thread: " + i);  
        }  
    }  
}
```

Implementing Runnable

```
public class RunnableExample implements Runnable {  
    public void run () {  
        for (int i = 1; i <= 100; i++) {  
            System.out.println ("Runnable: " + i);  
        }  
    }  
}
```

A Runnable Object

- The Thread object's **run()** method calls the Runnable object's **run()** method
- Allows threads to run inside any object, regardless of inheritance

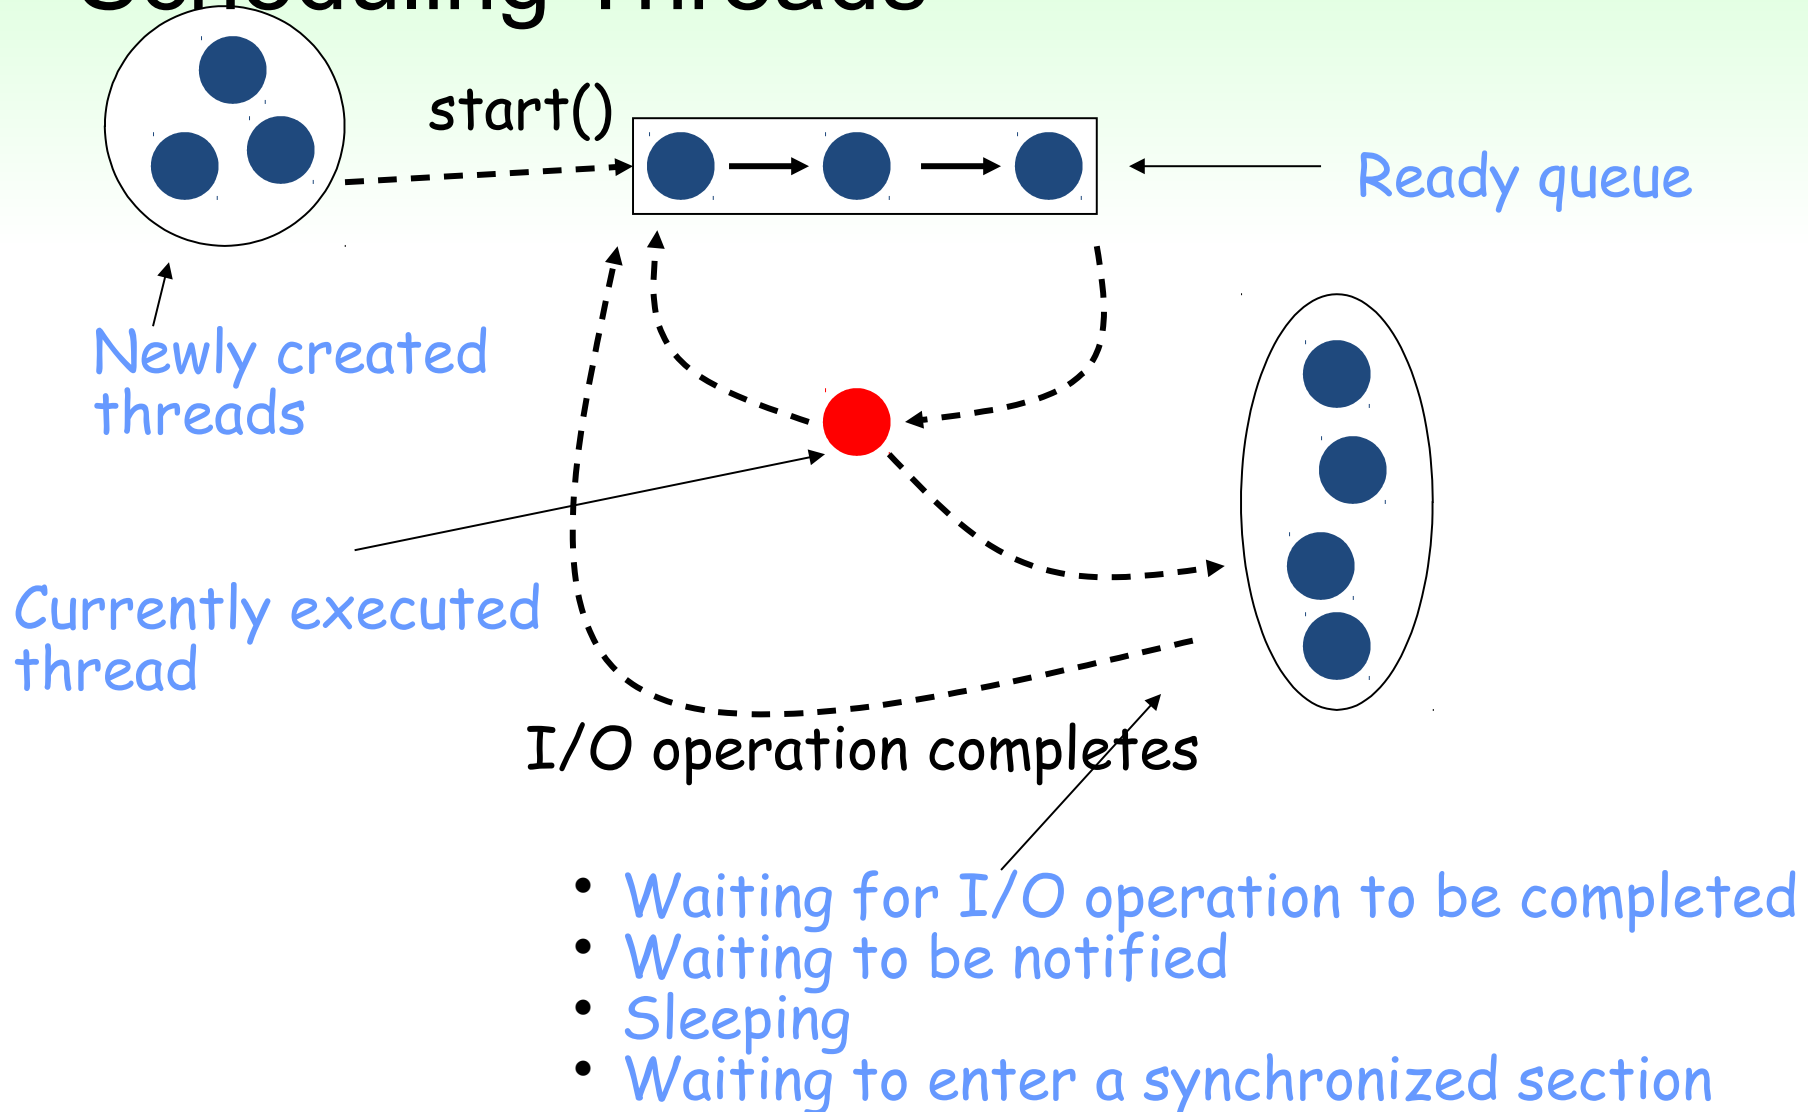
Example - a servlet that is also a thread

Starting the Threads

```
public class ThreadsStartExample {  
    public static void main (String argv[]) {  
        new ThreadExample ().start ();  
        new Thread(new RunnableExample ().start ());  
    }  
}
```

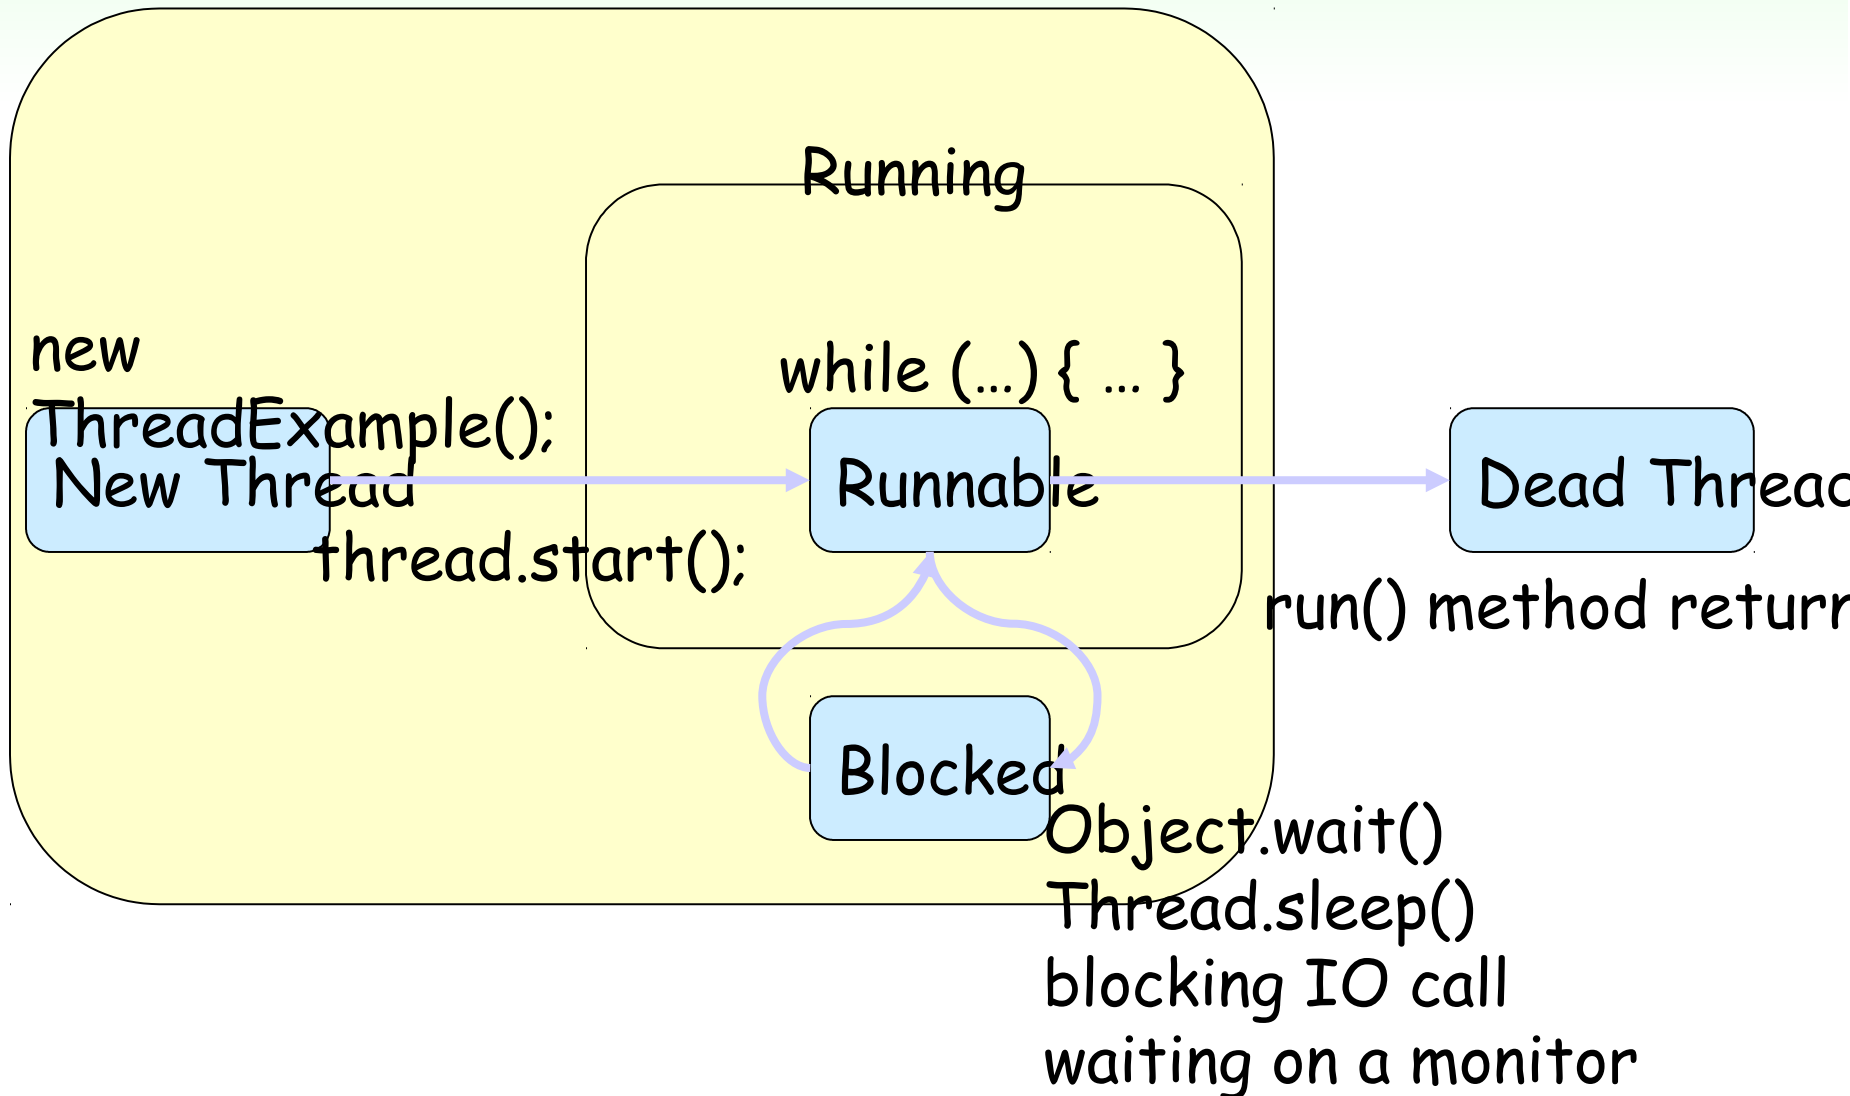
RESULT

Scheduling Threads



Thread State Diagram

Alive



Example

```
public class PrintThread1 extends Thread {
```

```
    String name;
```

```
    public PrintThread1(String name) {
```

```
        this.name = name;
```

```
    }
```

```
    public void run() {
```

```
        for (int i=1; i<500 ; i++) {
```

```
            try {
```

```
                sleep((long)(Math.random() * 100));
```

```
            } catch (InterruptedException ie) { }
```

```
            System.out.print(name);
```

```
        } }
```

Example (cont)

```
public static void main(String args[]) {  
    PrintThread1 a = new PrintThread1("*");  
    PrintThread1 b = new PrintThread1("-");  
    PrintThread1 c = new PrintThread1("=");  
  
    a.start();  
    b.start();  
    c.start();  
}  
}
```

RESULT

Scheduling

- Thread scheduling is the mechanism used to determine how runnable threads are allocated CPU time
- A thread-scheduling mechanism is either preemptive or nonpreemptive

Preemptive Scheduling

- Preemptive scheduling – the thread scheduler preempts (pauses) a running thread to allow different threads to execute
- Nonpreemptive scheduling – the scheduler never interrupts a running thread
- The nonpreemptive scheduler relies on the running thread to yield control of the CPU so that other threads may execute

Starvation

- A nonpreemptive scheduler may cause starvation (runnable threads, ready to be executed, wait to be executed in the CPU a lot of time, maybe even forever)
- Sometimes, starvation is also called a livelock

Time-Sliced Scheduling

- Time-sliced scheduling – the scheduler allocates a period of time that each thread can use the CPU
 - when that amount of time has elapsed, the scheduler preempts the thread and switches to a different thread
- Nontime-sliced scheduler – the scheduler does not use elapsed time to determine when to preempt a thread
 - it uses other criteria such as priority or I/O status

Java Scheduling

- Scheduler is preemptive and based on priority of threads
- Uses fixed-priority scheduling:
 - Threads are scheduled according to their priority w.r.t. other threads in the ready queue

Java Scheduling

- The highest priority runnable thread is always selected for execution above lower priority threads
- When multiple threads have equally high priorities, only one of those threads is guaranteed to be executing
- Java threads are guaranteed to be preemptive-but not time sliced
- **Q:** Why can't we guarantee time-sliced scheduling?

Thread Priority

- Every thread has a priority
- When a thread is created, it inherits the priority of the thread that created it
- The priority values range from 1 to 10, in increasing priority

Thread Priority (cont.)

- The priority can be adjusted subsequently using the **setPriority()** method
- The priority of a thread may be obtained using **getPriority()**
- Priority constants are defined:
 - MIN_PRIORITY=1
 - MAX_PRIORITY=10
 - NORM_PRIORITY=5

Some Notes

- Thread implementation in Java is actually based on operating system support
- Some Windows operating systems support only 7 priority levels, so different levels in Java may actually be mapped to the same operating system level
- What should we do about this?

Daemon Threads

- Daemon threads are “background” threads, that provide services to other threads, e.g., the garbage collection thread
- The Java VM will not exit if non-Daemon threads are executing
- The Java VM will exit if only Daemon threads are executing
- Daemon threads die when the Java VM exits


Concurrency

- An object in a program can be changed by more than one thread
- Q: Is the order of changes that were performed on the object important?

Race Condition

- A race condition – the outcome of a program is affected by the order in which the program's threads are allocated CPU time
- Two threads are simultaneously modifying a single object
- Both threads “race” to store their value

Race Condition Example



Put green pieces

How can we have
alternating colors?

Put red pieces

Monitors

- Each object has a “monitor” that is a token used to determine which application thread has control of a particular object instance
- In execution of a synchronized method (or block), access to the object monitor must be gained before the execution
- Access to the object monitor is queued

Monitor (cont.)

- Entering a monitor is also referred to as locking the monitor, or acquiring ownership of the monitor
- If a thread A tries to acquire ownership of a monitor and a different thread has already entered the monitor, the current thread (A) must wait until the other thread leaves the monitor

Critical Section

- The synchronized methods define critical sections
- Execution of critical sections is mutually exclusive. Why?

Example


```
public class BankAccount {
```

```
    private float balance;
```

```
    public synchronized void deposit(float amount) {  
        balance += amount;  
    }
```

```
    public synchronized void withdraw(float amount) {  
        balance -= amount;  
    }  
}
```

Critical Sections



Static Synchronized Methods

- Marking a static method as synchronized, associates a monitor with the class itself
- The execution of synchronized static methods of the same class is mutually exclusive. Why?

Example

```
public class PrintThread2 extends Thread {
```

```
    String name;
```

```
    public PrintThread2(String name) {
```

```
        this.name = name;
```

```
    }
```

```
    public static synchronized void print(String name) {
```

```
        for (int i=1; i<500 ; i++) {
```

```
            try {
```

```
                Thread.sleep((long)(Math.random() * 100));
```

```
            } catch (InterruptedException ie) { }
```

```
            System.out.print(str);
```

```
        }
```

```
    }
```

Example (cont)

```
public void run() {  
    print(name);  
}
```

```
public static void main(String args[]) {
```

```
    PrintThread2 a = new PrintThread2("*");
```

```
    PrintThread2 b = new PrintThread2("-");
```

```
    PrintThread2 c = new PrintThread2("=");
```

```
    a.start();
```

```
    b.start();
```

```
    c.start();
```

```
}
```

```
}
```

RESULT

```
public class BankAccount {
```

```
    private float balance;
```

```
    public synchronized void deposit(float amount) {  
        balance += amount;  
    }
```

```
    public synchronized void withdraw(float amount) {  
        balance -= amount;  
    }
```

```
    public synchronized void transfer(float amount,  
                                       BankAccount target) {  
        withdraw(amount);  
        target.deposit(amount);  
    }
```

```
}
```



```
public class MoneyTransfer implements Runnable {
```

```
    private BankAccount from, to;  
    private float amount;
```

```
    public MoneyTransfer(  
        BankAccount from, BankAccount to, float amount) {  
        this.from = from;  
        this.to = to;  
        this.amount = amount;  
    }
```

```
    public void run() {  
        from.transfer(amount, target);  
    }
```

```
}
```


```
BankAccount sureshAccount = new BankAccount();  
BankAccount rameshAccount = new BankAccount();
```

...

```
// At one place  
Runnable transaction1 =  
    new MoneyTransfer(sureshAccount, rameshAccount, 1200);  
Thread t1 = new Thread(transaction1);  
t1.start();
```

```
// At another place  
Runnable transaction2 =  
    new MoneyTransfer(rameshAccount, sureshAccount, 700);  
Thread t2 = new Thread(transaction2);  
t2.start();
```

Deadlocks



Java Locks are Reentrant

- Is there a problem with the following code?

```
public class Test {  
    public synchronized void a() {  
        b();  
        System.out.println("I am at a");  
    }  
    public synchronized void b() {  
        System.out.println("I am at b");  
    }  
}
```

Synchronized Statements

- A monitor can be assigned to a block
- It can be used to monitor access to a data element that is not an object, e.g., array
- Example:

```
void arrayShift(byte[] array, int count) {  
    synchronized(array) {  
        System.arraycopy (array, count,array, 0,  
            array.size - count);  
    }  
}
```

Thread Synchronization

- We need to synchronized between transactions, for example, the consumer-producer scenario



Wait and Notify

- Allows two threads to cooperate
- Based on a single shared lock object
 - Ramesh put a cookie wait and notify Suresh
 - Suresh eat a cookie wait and notify Ramesh
 - Ramesh put a cookie wait and notify Suresh
 - Suresh eat a cookie wait and notify Ramesh

The wait() Method

- The **wait()** method is part of the **java.lang.Object**
- It requires a lock on the object's monitor to execute
- It must be called from a synchronized method, or from a synchronized segment of code. Why?

The wait() Method

- wait() causes the current thread to wait until another thread invokes the notify() method or the notifyAll() method for this object
- Upon call for wait(), the thread releases ownership of this monitor and waits until another thread notifies the waiting threads of the object

The wait() Method

- **wait()** is also similar to **yield()**
 - Both take the current thread off the execution stack and force it to be rescheduled
- However, **wait()** is not automatically put back into the scheduler queue
 - **notify()** must be called in order to get a thread back into the scheduler's queue

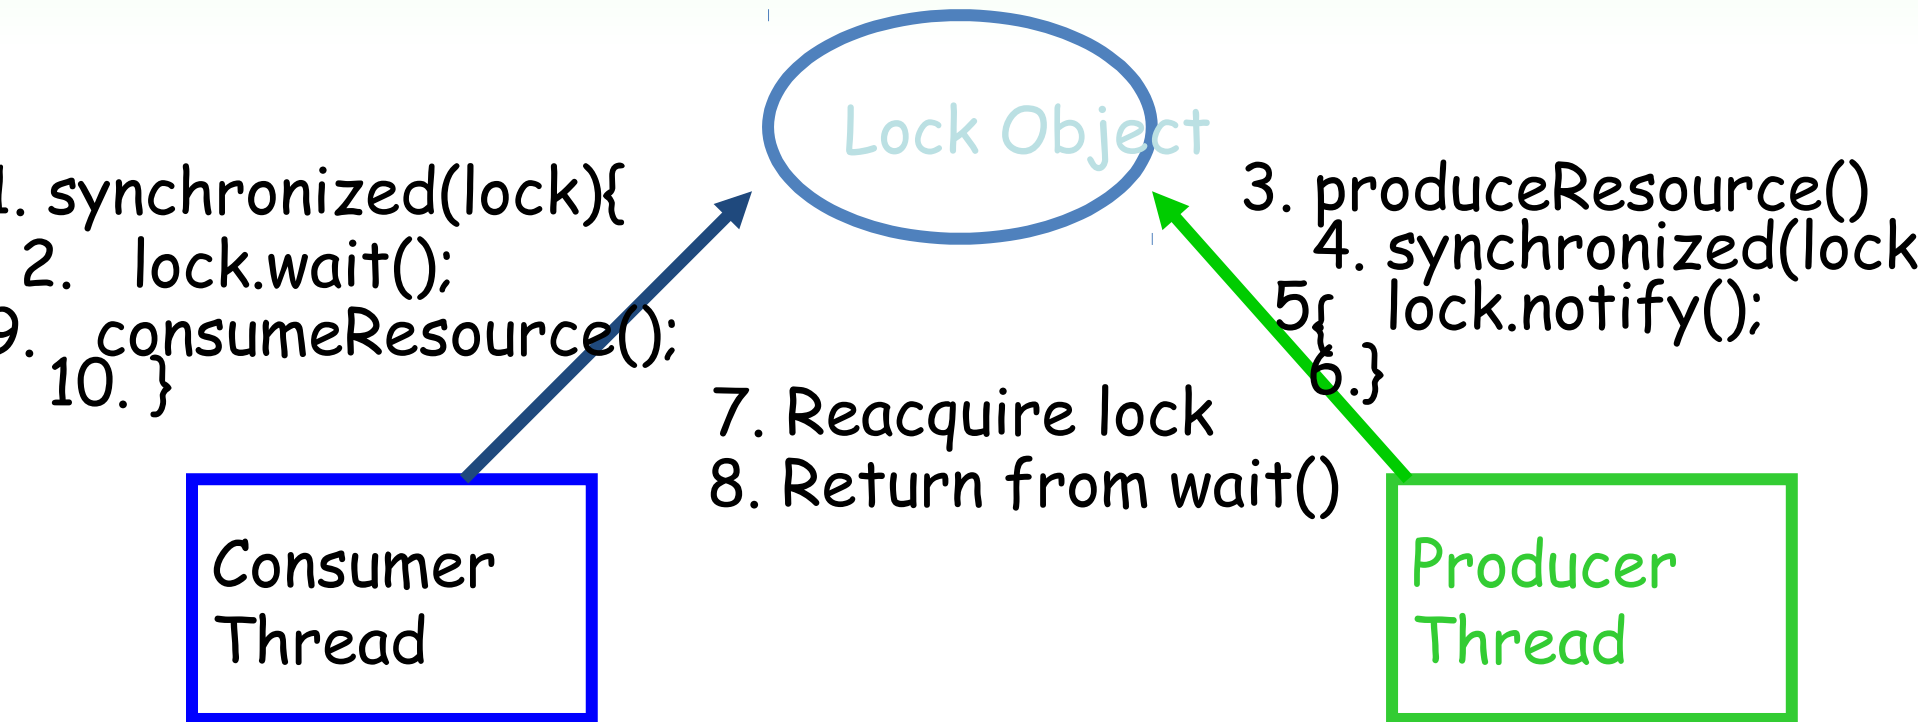
Consumer

```
synchronized (lock) {  
    while (!resourceAvailable()) {  
        lock.wait();  
    }  
    consumeResource();  
}
```

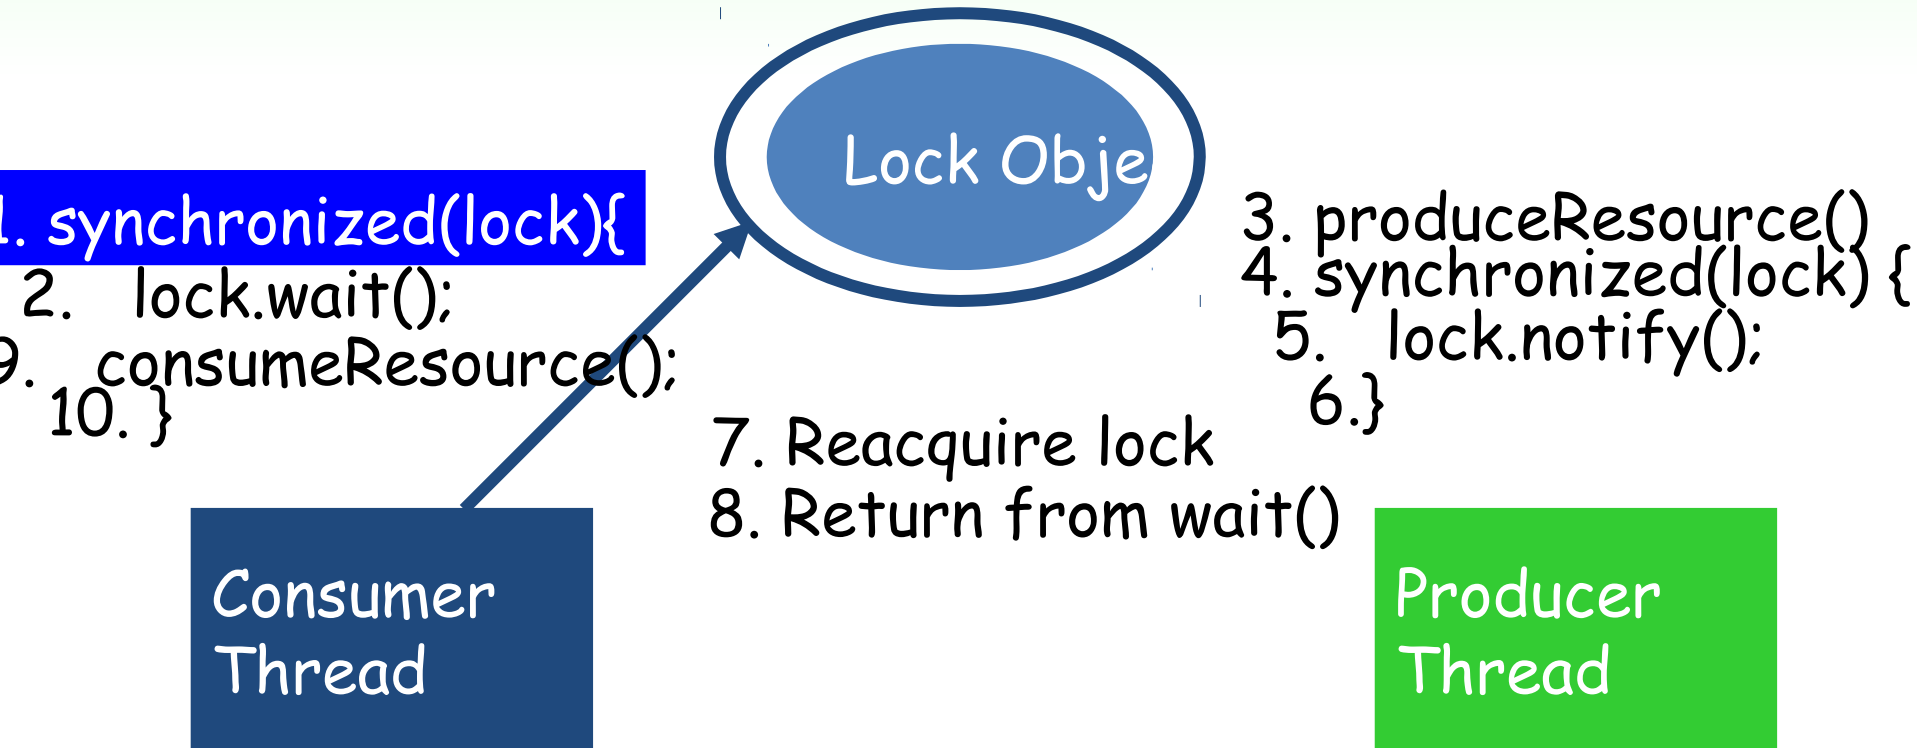
Producer

```
produceResource();  
synchronized (lock) {  
    lock.notifyAll();  
}
```

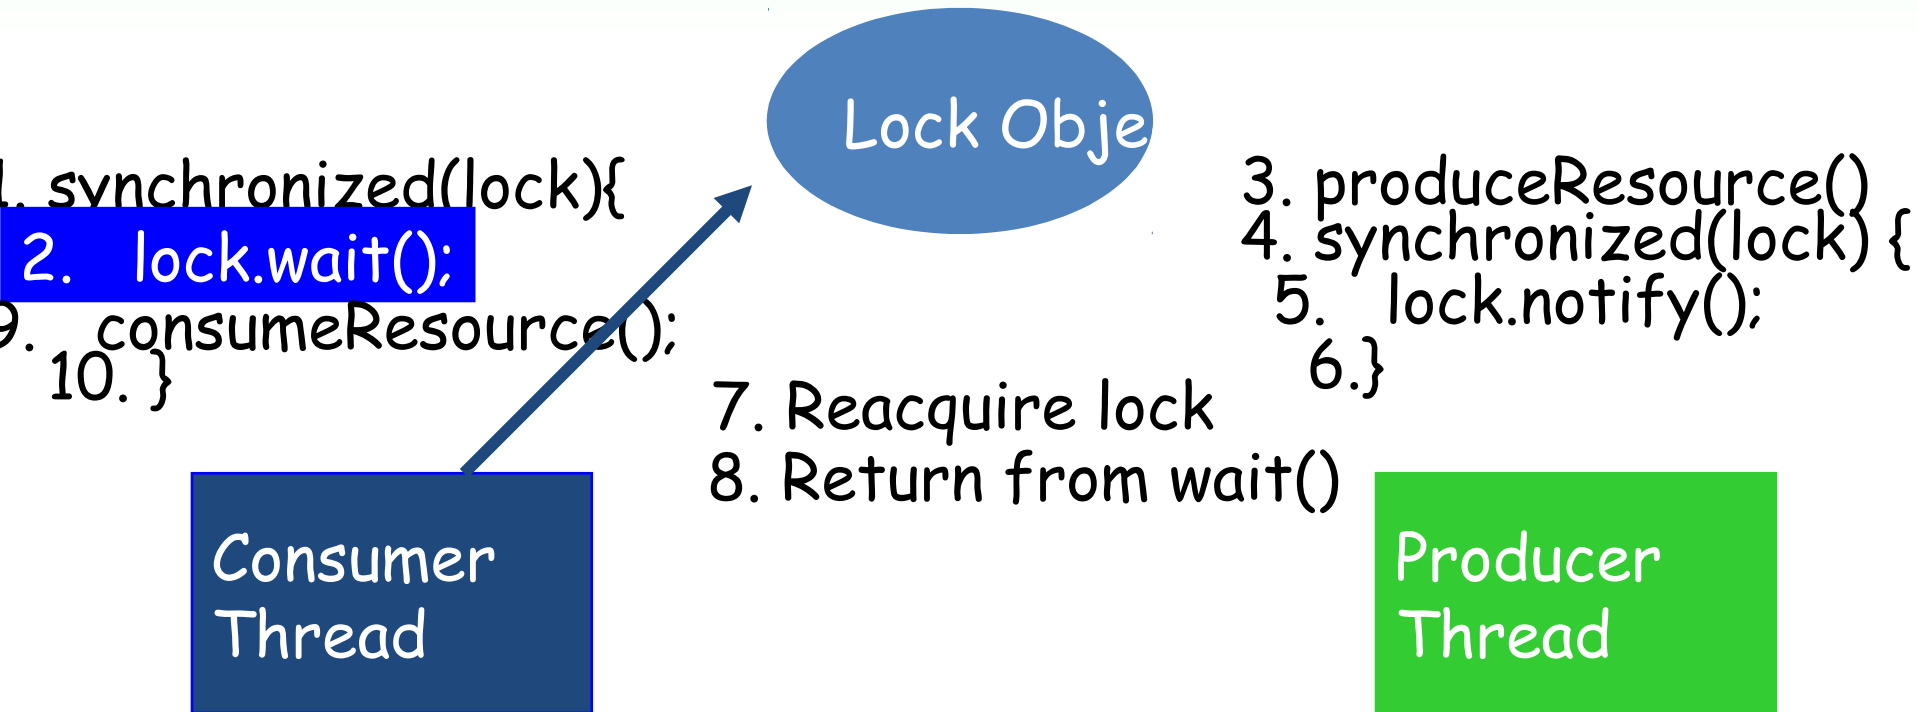
Wait/Notify Sequence



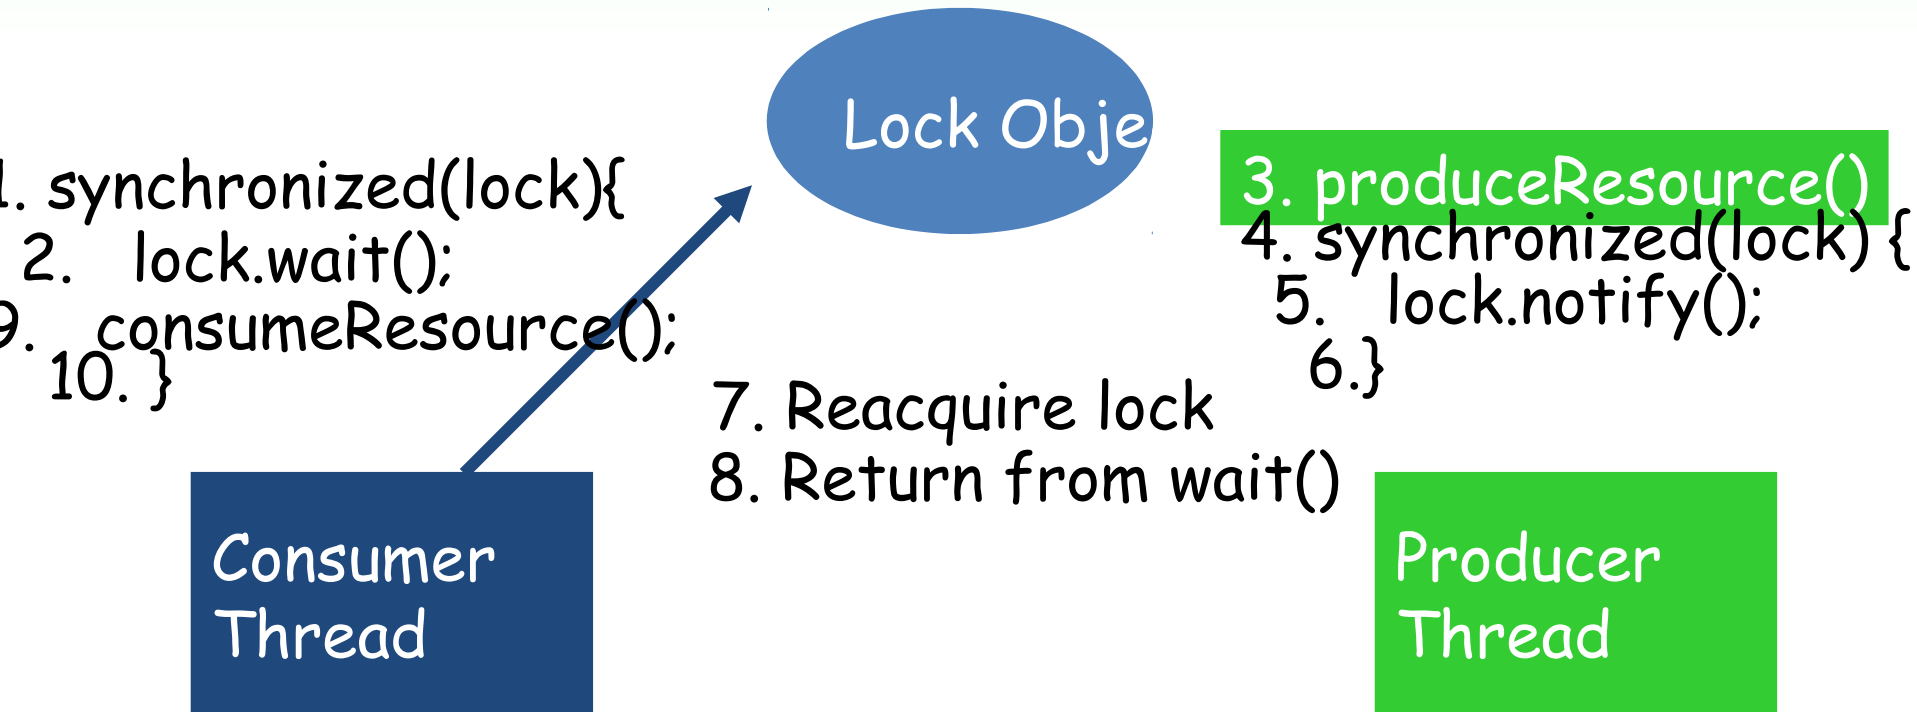
Wait/Notify Sequence



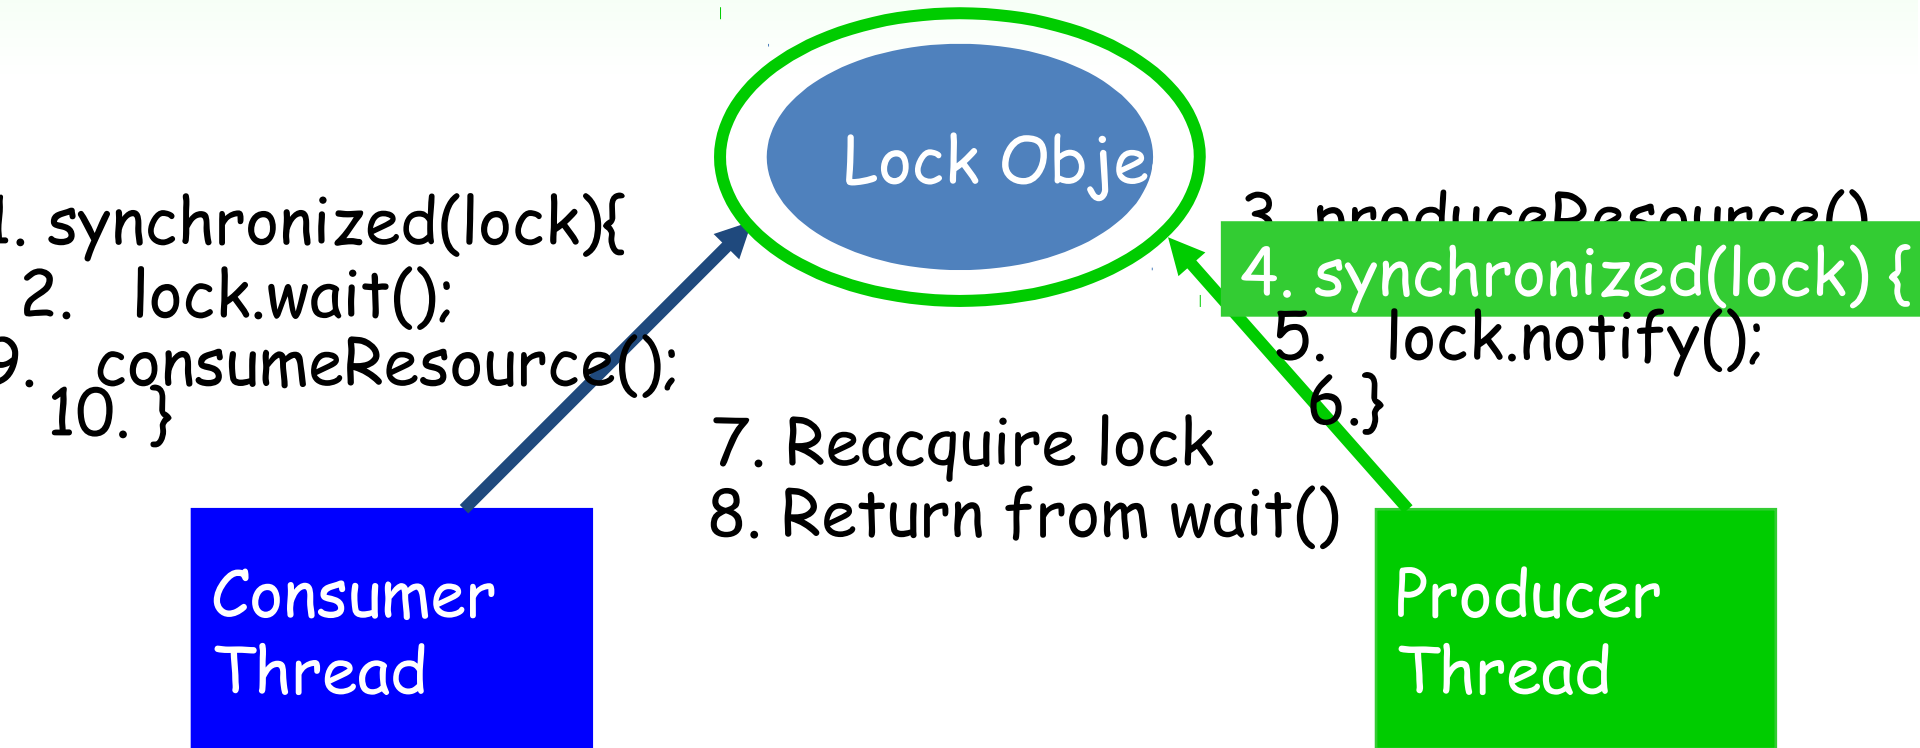
Wait/Notify Sequence



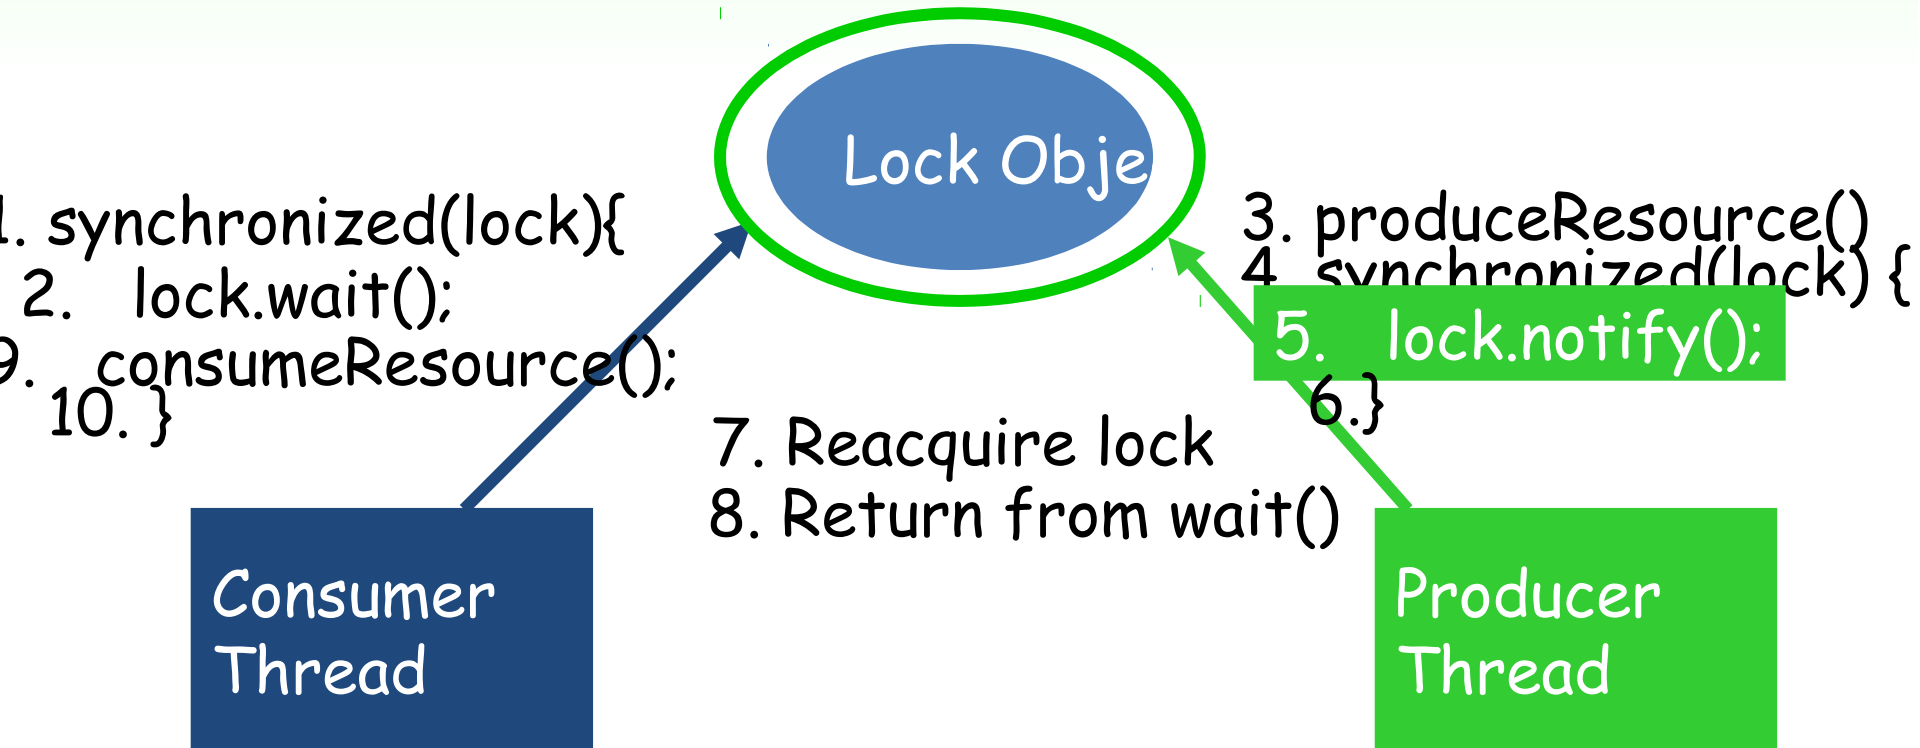
Wait/Notify Sequence



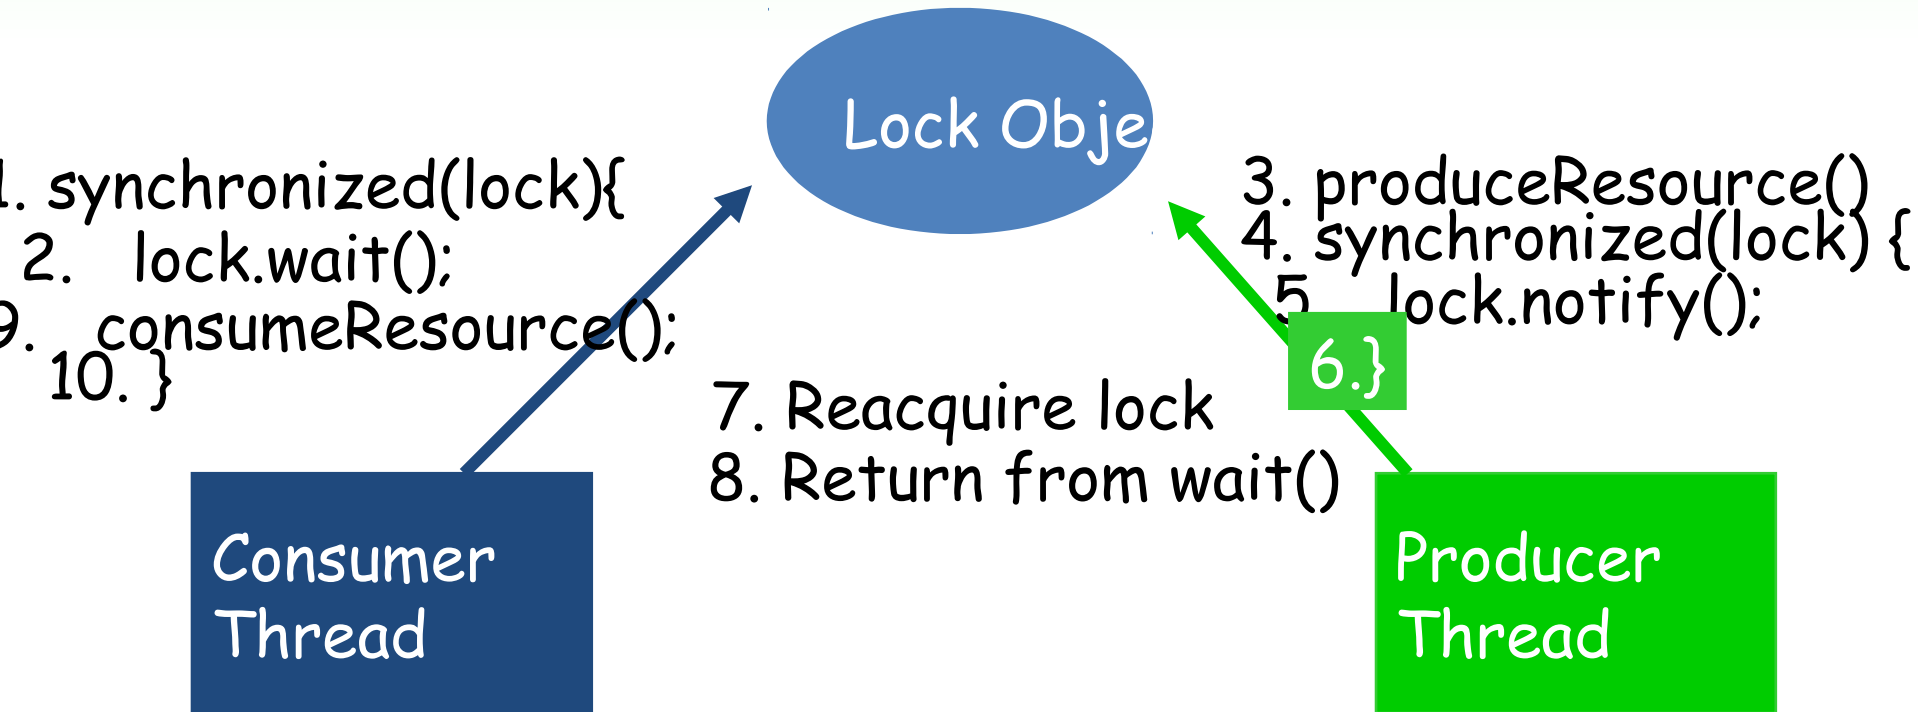
Wait/Notify Sequence



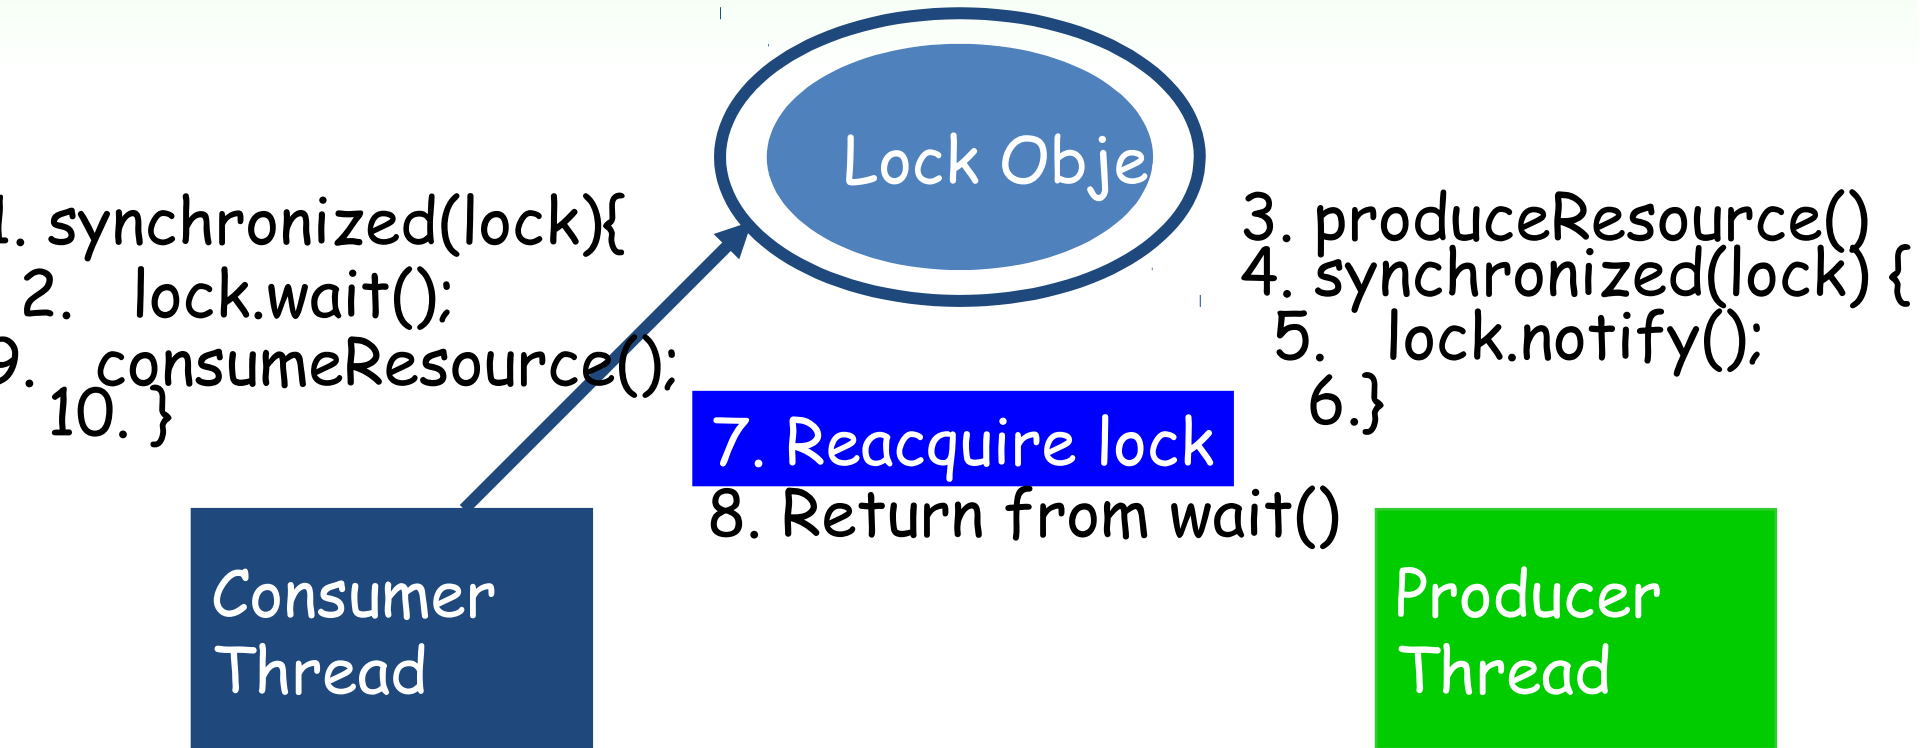
Wait/Notify Sequence



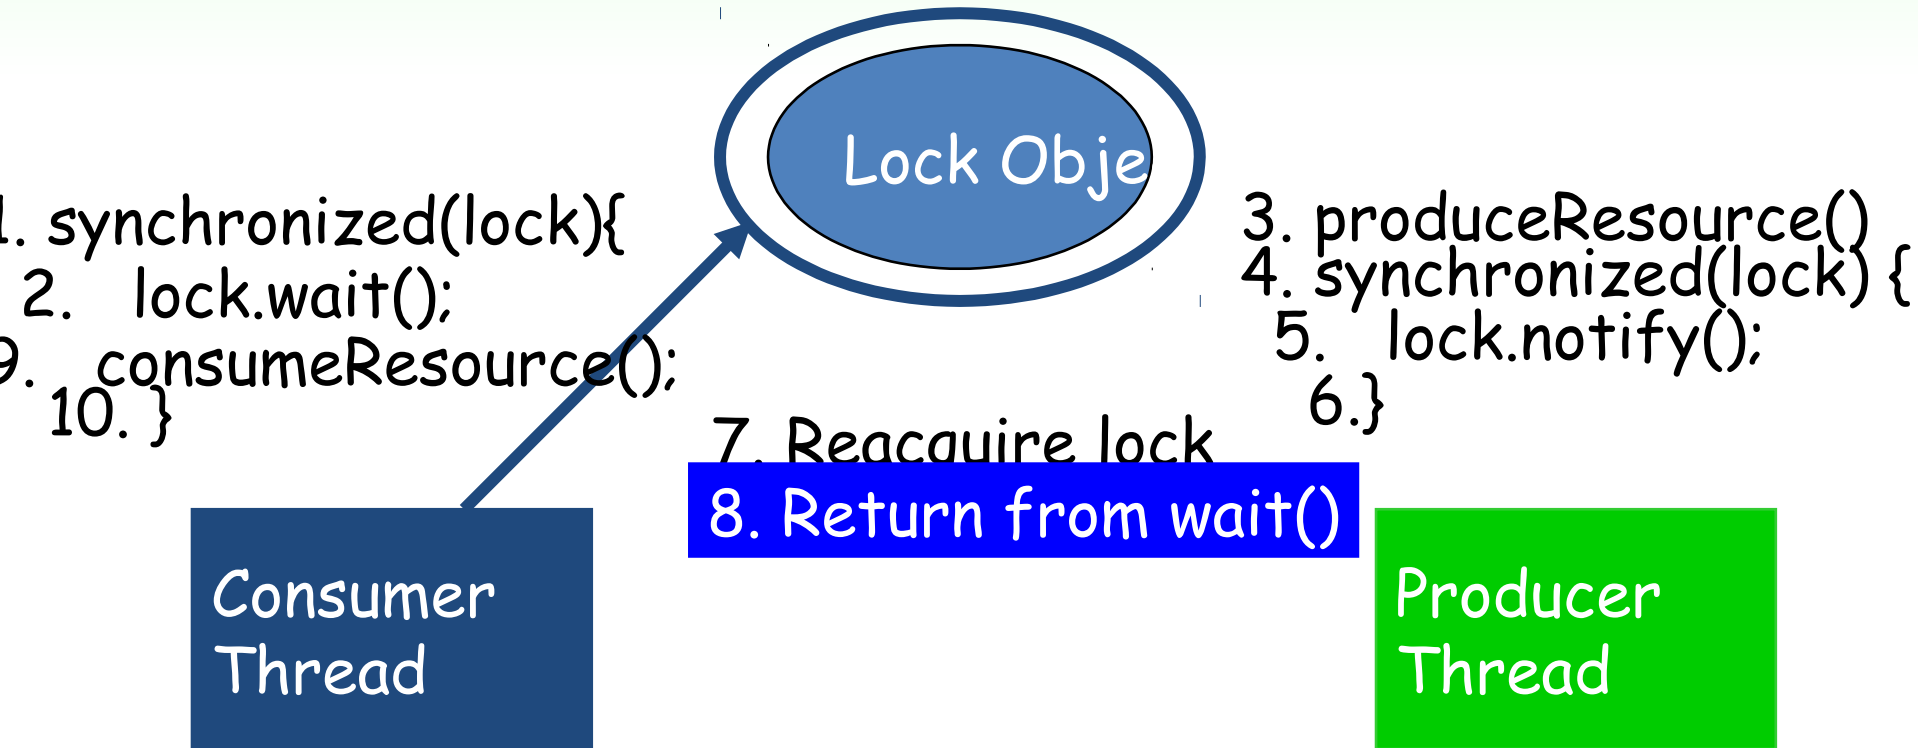
Wait/Notify Sequence



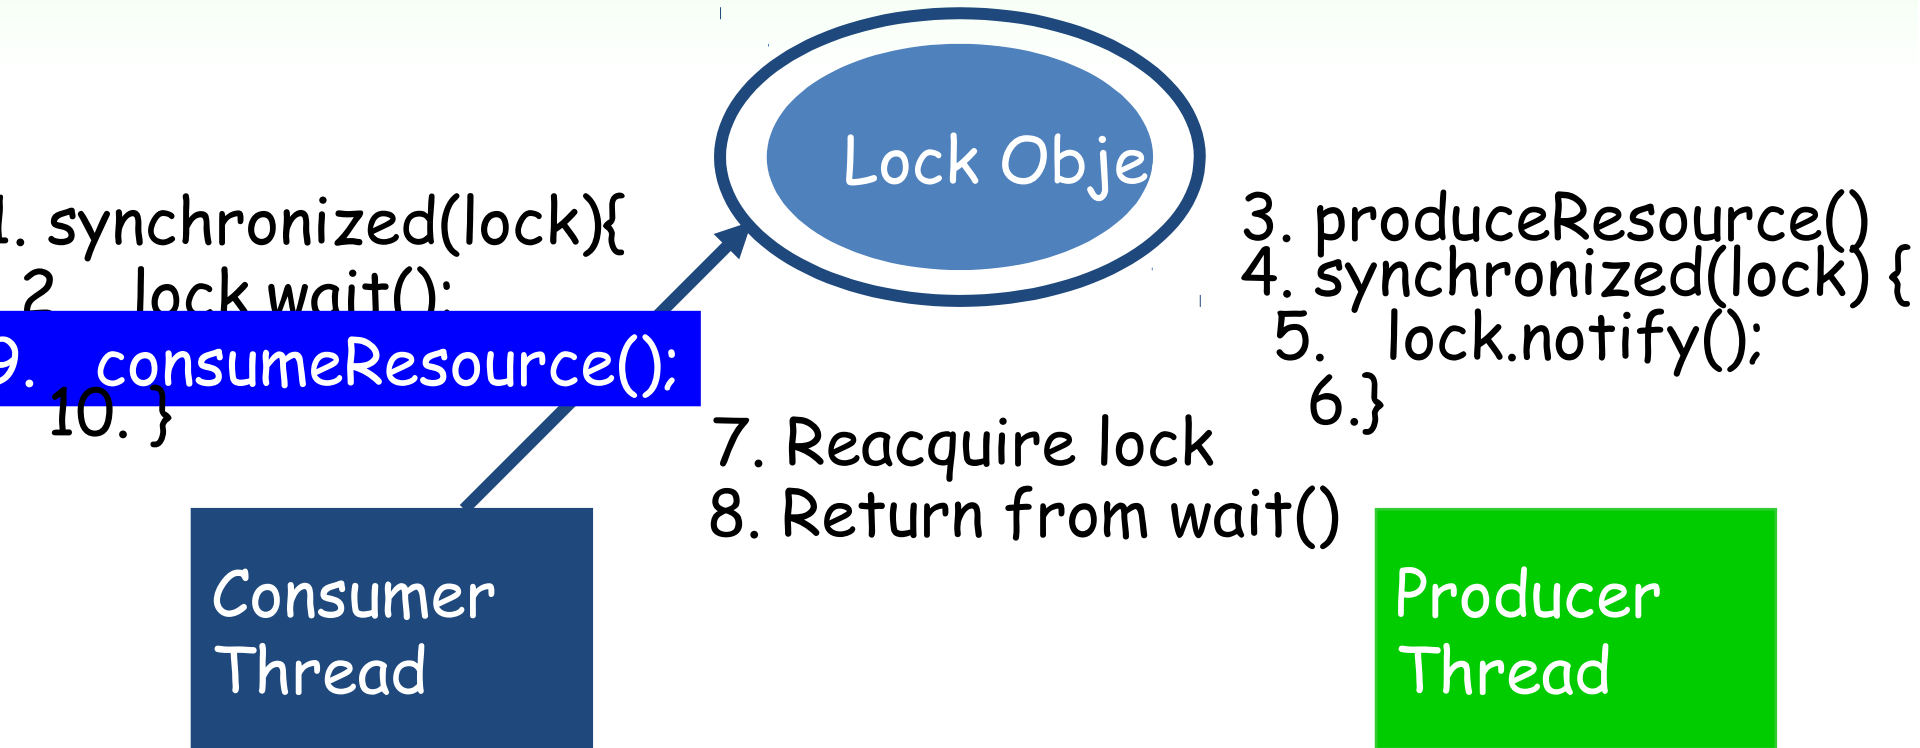
Wait/Notify Sequence



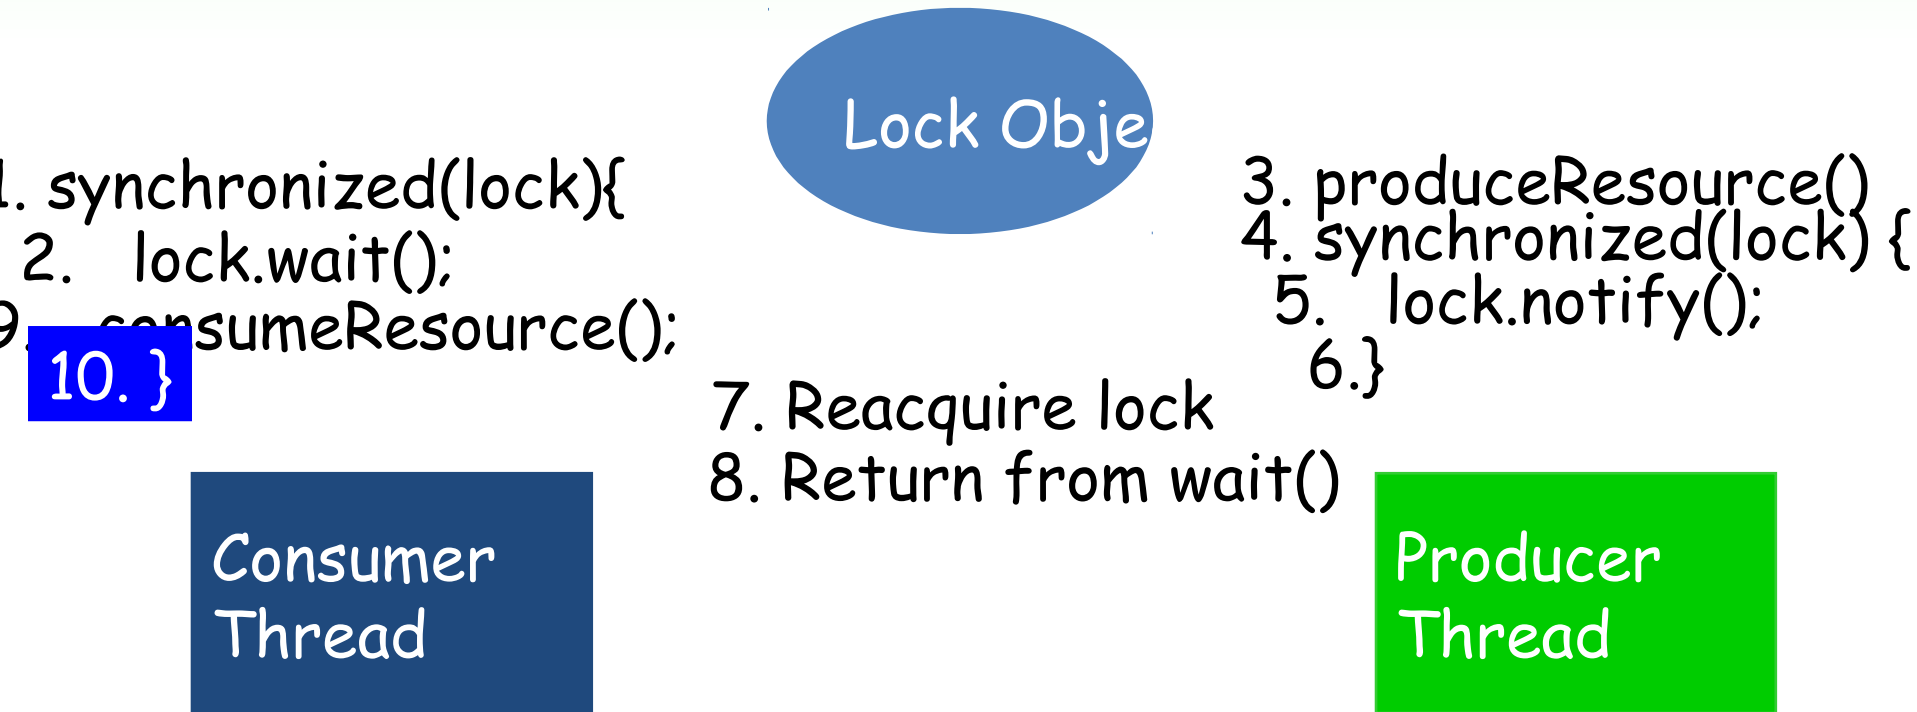
Wait/Notify Sequence



Wait/Notify Sequence



Wait/Notify Sequence



The Producer-Consumer Scenario: Cookies

```
public class Cookies {
```

```
    public static void main(String[] args) {
```

```
        CookieJar jar = new CookieJar();
```

```
        Suresh suresh = new Suresh(jar);  
        Ramesh ramesh = new Ramesh(jar);
```

```
        new Thread(Suresh).start();  
        new Thread(Ramesh).start();
```

```
    }
```

```
}
```


The Producer-Consumer

Scenario: Suresh

```
public class Suresh implements Runnable {
```

```
    CookyJar jar;
```

```
    public Suresh(CookyJar jar) {  
        this.jar = jar;  
    }
```

```
    public void eat() {  
        jar.getCooky("Suresh");  
        try {  
            Thread.sleep((int)Math.random() * 1000);  
        } catch (InterruptedException ie) {}  
    }
```

```
    public void run() {  
        for (int i = 1 ; i <= 10 ; i++) eat();  
    }  
}
```

The Producer-Consumer

Scenario: Ramesh

```
public class Ramesh implements Runnable {
```

```
    CookyJar jar;
```

```
    public Ramesh(CookyJar jar) {  
        this.jar = jar;  
    }
```

```
    public void bake(int cookyNumber) {  
        jar.putCooky("Ramesh", cookyNumber);  
        try {  
            Thread.sleep((int)Math.random() * 500);  
        } catch (InterruptedException ie) {}  
    }
```

```
    public void run() {  
        for (int i = 0 ; i < 10 ; i++) bake(i);  
    }  
}
```

CookieJar

```
public class CookyJar {  
    private int contents;  
    private boolean available = false;  
  
    public synchronized void getCooky(String who) {  
        while (!available) {  
            try {  
                wait();  
            } catch (InterruptedException e) { }  
        }  
        available = false;  
        notifyAll();  
        System.out.println( who + " ate cooky " + contents);  
    }  
}
```

CookieJar

```
public synchronized void putCooky(String who, int value) {  
    while (available) {  
        try {  
            wait();  
        } catch (InterruptedException e) { }  
    }  
    contents = value;  
    available = true;  
    System.out.println(who + " put cooky " + contents +  
                        " in the jar");  
    notifyAll();  
}}
```